

A Framework for Intent-Centric Computation

Abhimanyu Kashyap

aby@helloaxes.com

Abstract

The Aintent Virtual Machine (AVM) introduces a novel execution environment designed specifically for intent-centric computation in decentralized environments. Unlike traditional smart contract platforms that rely on deterministic function execution and manual orchestration, AVM natively interprets, decomposes, and executes high-level user intents through a verifiable, decentralized architecture. AVM's core innovation is the concept of "mutations"—intent-resolved state transitions—validated through zero-knowledge proofs and finalized via a novel consensus protocol called Proof of Mutations (PoM). This paper presents the formal design principles, execution semantics, cryptographic foundations, and consensus logic of AVM, providing a scalable and secure architecture for agentic systems, AI-coordinated workflows, and autonomous decentralized applications with quantum-resistant security considerations.

1. Introduction

The emergence of autonomous agents and AI-first interfaces has created a pressing need for infrastructure capable of understanding user goals, coordinating complex workflows, and verifying execution outcomes in decentralized environments. Current blockchain architectures and smart contract platforms are fundamentally limited in this emerging paradigm, as they require developers to explicitly encode logic, define deterministic triggers, and rely on economic incentives for correctness guarantees.

The Aintent Virtual Machine proposes a paradigm shift from function-level computation to intent-driven execution, where:

1. User goals are expressed in natural language or structured formats
2. These expressions are resolved into formalized execution plans
3. Execution is decomposed into atomic, verifiable operations
4. Each state transition is cryptographically proven and validated
5. Results are committed to a global state with instant finality

AVM introduces Deep Intents, DeepFlows, Quarks, and Mutations as the primitive constructs that define how workflows are parsed, orchestrated, verified, and committed to the global state. The system operates in conjunction with the Proof of Mutations consensus mechanism,

enabling provable execution integrity, instant finality, and zero-knowledge based accountability at every step of state transition.

1.1 Comparison with Existing Approaches

Feature	Traditional Smart Contracts	Aintent Virtual Machine
Input Format	Function calls with explicit parameters	Natural language intents or structured goals
Execution Model	Deterministic, sequential	Goal-oriented, parallelizable
State Transitions	Direct state modifications	Proven mutations with zk-verification
Finality	Probabilistic (in PoW/PoS)	Deterministic (proof-based)
Composability	Manual orchestration	Automatic decomposition and execution
Security Model	Economic (stake/cost)	Cryptographic (zk-proofs)

2. Execution Model Overview

In AVM, every user-submitted intent undergoes a structured pipeline from expression to execution:

- Intent Parsing:** Transformer-based language models convert natural language input into structured Directed Acyclic Graphs (DAGs) representing the goal and its subcomponents
- Workflow Decomposition:** These DAGs are analyzed and split into atomic executable steps called Quarks
- Mutation Generation:** Each Quark executes its designated function and generates a proposed state delta with a corresponding zero-knowledge proof
- Proof Verification:** Validators cryptographically verify that all mutations are correct and non-conflicting
- State Finalization:** Valid mutations are appended to the canonical chain with deterministic finality

Each mutation μ is defined as a verifiable execution unit comprising:

- A cryptographic reference to the Deep Intent
- Input and output state deltas (δ)
- A zero-knowledge proof (π) attesting to correct computation
- A cryptographic signature binding the execution to its prover

3. Key Components

3.1 Deep Intent

A Deep Intent I is a semantic abstraction representing a user's goal, formalized as a directed acyclic graph (DAG). Mathematically, it is defined as:

$$I = (G, \Phi, C)$$

Where:

- $G = \{g_1, g_2, \dots, g_n\}$ is the set of subgoals
- $\Phi = \{\phi_1, \phi_2, \dots, \phi_m\}$ represents the dependencies between subgoals
- $C = \{c_1, c_2, \dots, c_p\}$ defines constraints on execution (e.g., timing, resource limits)

A Deep Intent may yield multiple candidate DeepFlows, each representing an executable path through the DAG based on resource constraints, user preferences, or system state.

3.2 DeepFlow

A DeepFlow D is a concrete execution graph derived from a Deep Intent. Formally:

$$D = (Q, E)$$

Where:

- $Q = \{q_1, q_2, \dots, q_k\}$ is the set of Quarks (atomic execution units)
- $E = \{e_1, e_2, \dots, e_j\}$ represents edges defining data flow or execution dependencies

Each DeepFlow is optimized for latency, cost, and verifiability through graph transformation algorithms.

3.3 Quark

A Quark q is the atomic execution unit in AVM, responsible for resolving a single subgoal within a DeepFlow. Each Quark q is defined by the tuple:

$$q = (f, \sigma_{in}, \sigma_{out}, \pi)$$

Where:

- f is the function to be executed
- σ_{in} is the input state
- σ_{out} is the output state after execution

- π is the zero-knowledge proof of correct computation

Each Quark performs three essential operations:

1. Reads relevant input state
2. Performs a bounded transformation
3. Emits an output and zero-knowledge proof of correct computation

3.4 Mutation

A mutation μ is a bundle of Quark outputs submitted for inclusion in the ledger. Formally:

$$\mu = (I, D, \Delta, \Pi, h)$$

Where:

- I is the reference to the original Deep Intent
- D is the DeepFlow identifier
- $\Delta = \{\delta_1, \delta_2, \dots, \delta_n\}$ is the set of state deltas
- $\Pi = \{\pi_1, \pi_2, \dots, \pi_n\}$ is the set of corresponding zk-proofs
- h is the commitment hash of the resulting state

3.5 Verifier Set

The verifier set V consists of validators responsible for:

1. Checking the validity of each zk-proof
2. Verifying non-conflict between concurrent mutations
3. Committing final state transitions in PoM blocks

4. Proof of Mutations (PoM) Consensus

The AVM chain is governed by the PoM consensus protocol, which prioritizes execution correctness over block proposer identity or stake weight.

4.1 Block Structure

A block B in the AVM blockchain is defined as:

$$B = (H, M, S, \sigma)$$

Where:

- H is the block header containing metadata

- $M = \{\mu_1, \mu_2, \dots, \mu_r\}$ is the set of mutations included in the block
- S is the state root after applying all mutations
- σ is the aggregate signature of validators

4.2 Validity Conditions

A block B is considered valid if and only if all contained mutations $\mu_i \in M$:

1. Pass zk-proof verification: $V(\pi_{i,j}, \delta_{i,j}, S) = \text{true}$ for all proofs $\pi_{i,j} \in \mu_i$
2. Do not conflict with each other: $\forall \mu_i, \mu_j \in M, i \neq j: \text{Conflict}(\mu_i, \mu_j) = \text{false}$
3. Result in a valid next-state hash: $S' = \text{Apply}(M, S)$ and $\text{Hash}(S')$ matches the committed state root

The PoM consensus guarantees deterministic finality, low latency, and scalable parallel verification.

4.3 Consensus Algorithm

Algorithm 1: PoM Consensus Protocol

Input: Set of pending mutations M_{pending}

Output: Valid block B or rejection

```

1. function ProposeMutations( $M_{\text{pending}}$ )
2.    $M_{\text{selected}} \leftarrow \emptyset$ 
3.   for each  $\mu \in M_{\text{pending}}$  do
4.      $\text{valid} \leftarrow \text{true}$ 
5.     for each  $\pi \in \mu.\Pi$  do
6.       if !VerifyProof( $\pi, \mu.\delta, S$ ) then
7.          $\text{valid} \leftarrow \text{false}$ 
8.         break
9.       end if
10.    end for
11.    if  $\text{valid}$  and !ConflictsWith( $\mu, M_{\text{selected}}$ ) then
12.       $M_{\text{selected}} \leftarrow M_{\text{selected}} \cup \{\mu\}$ 
13.    end if
14.  end for
15.   $S' \leftarrow \text{Apply}(M_{\text{selected}}, S)$ 
16.   $B \leftarrow \text{CreateBlock}(M_{\text{selected}}, \text{Hash}(S'))$ 
17.  return  $B$ 
18. end function

```

```

19. function ValidationPhase(B)
20.   valid ← true
21.   for each  $\mu \in B.M$  do
22.     for each  $\pi \in \mu.\Pi$  do
23.       if !VerifyProof( $\pi$ ,  $\mu.\delta$ , S) then
24.         valid ← false
25.         break
26.       end if
27.     end for
28.     if !valid then break end if
29.   end for
30.   if valid and Hash(Apply(B.M, S)) = B.S then
31.     SignAndBroadcast(B)
32.     return true
33.   else
34.     RejectBlock(B)
35.     return false
36.   end if
37. end function

```

5. Formal Execution Semantics

Let I be a user-submitted intent. The AVM execution semantics are formally defined through the following stages:

5.1 Intent Decomposition

The Deep Intent I is parsed into a set of candidate DeepFlows $\{D_1, D_2, \dots, D_n\}$:

Algorithm 2: Intent Decomposition

Input: User intent expression I_expr

Output: Set of candidate DeepFlows $\{D_1, D_2, \dots, D_n\}$

```

1. function DecomposeIntent( $I\_expr$ )
2.    $I \leftarrow \text{ParseIntent}(I\_expr)$  // Convert to structured intent
3.    $G \leftarrow \text{ExtractSubgoals}(I)$  // Generate subgoal graph
4.    $D\_candidates \leftarrow \emptyset$ 
5.   for each strategy  $s \in \text{DecompositionStrategies}$  do
6.      $D\_s \leftarrow s.\text{ApplyTo}(G)$  // Apply decomposition strategy

```

```

7.         if IsValidDeepFlow(D_s) then
8.             D_candidates  $\leftarrow$  D_candidates  $\cup$  {D_s}
9.         end if
10.    end for
11.    return D_candidates
12. end function

```

Each DeepFlow D_i is a DAG defined as $D_i = (Q_i, E_i)$, where Q_i are Quark nodes and E_i are dependency edges.

5.2 DeepFlow Selection

The user or optimization engine selects a preferred DeepFlow $D^* \in \{D_i\}$ based on metrics like time, cost, trust, or resource availability:

Algorithm 3: DeepFlow Selection

Input: Set of candidate DeepFlows $\{D_1, D_2, \dots, D_n\}$, selection criteria C

Output: Selected DeepFlow D^*

```

1. function SelectDeepFlow(D_candidates, C)
2.     scores  $\leftarrow$   $\emptyset$ 
3.     for each D  $\in$  D_candidates do
4.         score  $\leftarrow$  EvaluateFlow(D, C) // Score based on criteria
5.         scores  $\leftarrow$  scores  $\cup$  {(D, score)}
6.     end for
7.     D*  $\leftarrow$  MaxScore(scores)
8.     return D*
9. end function

```

5.3 Quark Execution and Mutation Generation

Each node $q_j \in D^*$ is executed independently (where dependencies allow):

Algorithm 4: Quark Execution

Input: DeepFlow $D^* = (Q, E)$, global state S

Output: Set of mutations M

```

1. function ExecuteQuarks(D*, S)

```

```

2.   M ← ∅
3.   executed ← ∅
4.   ready ← {q ∈ Q | ∀e = (q', q) ∈ E: q' ∈ executed}
5.
6.   while ready ≠ ∅ do
7.       parallel for each q ∈ ready do
8.           σin ← ExtractInputState(q, S)
9.           (σout, π) ← ExecuteQuark(q, σin)
10.          δ ← ComputeDelta(σin, σout)
11.          executed ← executed ∪ {q}
12.          M ← M ∪ {(q, δ, π)}
13.       end parallel for
14.       ready ← {q ∈ Q | ∀e = (q', q) ∈ E: q' ∈ executed} \ executed
15.   end while
16.
17.   μ ← PackageMutation(D*, M)
18.   return μ
19. end function

```

For each executed Quark q_j , a delta δ_j and proof π_j are emitted. Mutations are constructed as $\mu = \{\delta_j, \pi_j, h_{\text{expected}}\}$.

5.4 Verification and Commitment

Verifiers ensure $V(\pi_j, \delta_j, S) = \text{true}$. Mutations that pass verification are appended to the chain in the next block B_t .

6. Security Properties and Fault Tolerance

6.1 Integrity of Execution

AVM ensures execution integrity through several mechanisms:

1. All execution steps must be accompanied by zero-knowledge proofs
2. Proofs are non-interactive and succinct, minimizing verifier load
3. Invalid or unverifiable mutations are rejected deterministically
4. Malicious Quarks proposing false state transitions cannot produce valid zk-proofs

The cryptographic security of AVM's execution model relies on the soundness properties of the underlying zero-knowledge proof system. We employ a hybrid approach using different zk-proof systems based on computation complexity:

- **SNARKs (Succinct Non-interactive Arguments of Knowledge)** for computationally intensive operations
- **STARKs (Scalable Transparent Arguments of Knowledge)** for quantum resistance
- **Bulletproofs** for range proofs and simple computations

6.2 Resistance to Byzantine Actors

AVM is designed to maintain security in the presence of Byzantine actors:

1. Malicious validators who falsely attest to invalid proofs can be challenged and slashed via AVM's dispute protocol
2. The protocol remains secure as long as the majority of verification power is honest
3. Security guarantees hold under the standard asynchronous Byzantine model assuming less than 1/3 of nodes are malicious

6.3 Post-Quantum Security

To ensure long-term security in the post-quantum era, AVM incorporates several quantum-resistant cryptographic primitives:

Algorithm 5: Quantum-Resistant Proof Generation

Input: Quark q , input state σ_{in} , output state σ_{out}

Output: Quantum-resistant proof π

```

1. function GenerateQuantumResistantProof( $q, \sigma_{in}, \sigma_{out}$ )
2.     // Hash using quantum-resistant hash function (SPHINCS+)
3.      $h_{in} \leftarrow \text{SPHINCS+}(\sigma_{in})$ 
4.      $h_{out} \leftarrow \text{SPHINCS+}(\sigma_{out})$ 
5.
6.     // Generate STARK proof (quantum-resistant)
7.      $\text{computation\_trace} \leftarrow \text{GenerateExecutionTrace}(q, \sigma_{in}, \sigma_{out})$ 
8.      $\pi_{stark} \leftarrow \text{GenerateSTARKProof}(\text{computation\_trace})$ 
9.
10.    // Sign with quantum-resistant signature (CRYSTALS-Dilithium)
11.     $\text{signature} \leftarrow \text{Dilithium-Sign}(h_{in} || h_{out} || \pi_{stark})$ 
12.
13.    return ( $\pi_{stark}, \text{signature}$ )
14. end function

```

We employ the following quantum-resistant cryptographic primitives:

1. **SPHINCS+** for hash-based signatures
2. **CRYSTALS-Dilithium** for lattice-based signatures
3. **STARKs** for zero-knowledge proofs (inherently quantum-resistant due to reliance on hash functions)
4. **CRYSTALS-Kyber** for key encapsulation when encryption is needed

6.4 Finality and Fork Resistance

Unlike probabilistic consensus mechanisms, PoM provides deterministic finality:

1. Since all valid mutations are proven and verified, finality is deterministic
2. A block containing a mutation becomes final as soon as all zk-proofs are verified and committed
3. This eliminates probabilistic finality risks present in PoW and PoS systems

6.5 Partition Tolerance

AVM is designed to handle network partitions gracefully:

1. The PoM model ensures that honest partitions maintain a canonical state independently
2. Conflicting states due to network splits are resolved through deterministic re-orgs using mutation history and proof validity
3. Partitioned networks can continue processing independent DeepFlows without affecting overall security

7. Performance, Scalability & Benchmarks

The Aintent Virtual Machine is architected to support high-throughput, low-latency intent execution through parallelized verification, zk-proof compression, and decentralized solver sets.

7.1 Parallel Quark Execution

Theoretical analysis of AVM's parallelism shows significant performance advantages over sequential execution models:

If k Quarks are independent in a DeepFlow, they can be executed and verified in $O(\log k)$ time via parallel scheduling and batch zk-verification, compared to $O(k)$ time in sequential models.

Algorithm 6: Parallel Quark Scheduler

Input: DeepFlow $D = (Q, E)$

Output: Execution schedule with maximum parallelism

```

1. function GenerateParallelSchedule(D)
2.   levels ← TopologicalSort(D)
3.   schedule ← []
4.
5.   for each level L ∈ levels do
6.     parallel_batch ← {q | q ∈ L}
7.     schedule.append(parallel_batch)
8.   end for
9.
10.  return schedule
11. end function

```

7.2 zk-Proof Aggregation

AVM supports recursive proof composition to compress multiple Quark proofs into a single verification step:

Algorithm 7: Recursive Proof Aggregation

Input: Set of proofs $\Pi = \{\pi_1, \pi_2, \dots, \pi_n\}$

Output: Aggregated proof $\pi_{a \circ \circ}$

```

1. function AggregateProofs( $\Pi$ )
2.   // Base case
3.   if  $|\Pi| = 1$  then
4.     return  $\Pi[0]$ 
5.   end if
6.
7.   // Recursive aggregation in balanced binary tree
8.   mid ←  $|\Pi| / 2$ 
9.   left_proofs ←  $\Pi[0:mid]$ 
10.  right_proofs ←  $\Pi[mid:|\Pi|]$ 
11.
12.   $\pi_{left}$  ← AggregateProofs(left_proofs)
13.   $\pi_{right}$  ← AggregateProofs(right_proofs)
14.
15.   $\pi_{agg}$  ← ProveAggregation( $\pi_{left}$ ,  $\pi_{right}$ )
16.  return  $\pi_{agg}$ 
17. end function

```

This results in improved verifier efficiency without compromising proof integrity or atomicity of execution.

7.3 Verifier Throughput

With optimization of the verification circuits and hardware acceleration, AVM achieves significant throughput:

- With succinct proofs (~1–2 KB) and native AVM verifier circuits, a consumer-grade validator node can verify 10,000+ Quarks per second
- With GPU acceleration, throughput can exceed 100,000 verified Quarks/sec
- Recursive SNARK verification enables constant-time verification regardless of computation complexity

7.4 Block Finality Time

Target finality is less than 5 seconds per block, even under adversarial network latency, due to:

1. Parallel proof verification
2. Efficient consensus protocol that doesn't require probabilistic confirmation
3. Optimized network messaging patterns for proof distribution

7.5 Storage Optimization

AVM employs sophisticated storage optimization techniques:

1. Only commitment hashes of state transitions and zk-proofs are stored on-chain
2. Full state deltas and execution transcripts are stored off-chain but retrievable via content-addressable storage (similar to IPFS)
3. The system employs a Merkle-based accumulator for efficient state retention and retrieval

8. Use Cases & Applications

The Aintent Virtual Machine provides a universal platform for verifiable, goal-oriented computation across multiple domains:

8.1 Autonomous Agents

AVM powers AI-native agents that translate user goals into executable, auditable workflows. For example, a travel-planning agent could book flights, hotels, and handle refunds automatically—with zk-proofs at every step.

This enables:

- Verifiable delegation to autonomous systems
- Composable agent behaviors with cryptographic guarantees
- Fault-tolerant execution with recovery capabilities

8.2 Decentralized Finance (DeFi)

AVM enables verifiable swaps, lending strategies, and automated liquidity management. For instance, an intent to "allocate \$500 across ETH and stablecoins" could be executed through multiple Quarks selecting the best pools and generating composable proofs.

Key advantages include:

- Intent-based trading rather than function-based interaction
- Automatic strategy optimization based on user goals
- Cryptographic guarantees of best execution

8.3 DAO Governance Workflows

Decentralized Autonomous Organizations can encode operational goals as Deep Intents, such as fund transfers or bounty creation. A DAO vote could trigger a multi-step intent to pay contributors, allocate budget, and publish receipts—each mutation audited and finalized on-chain.

This provides:

- Automated governance execution
- Cryptographic accountability for treasury operations
- Complex multi-stage workflows with verification

8.4 E-Commerce & Marketplace Automation

Buyers can issue intents like "Find me the cheapest verified option for X." Sellers and fulfillment agents (Quarks) can then compete and execute automatically, with each step verified.

Benefits include:

- Goal-oriented purchasing rather than vendor selection
- Competitive fulfillment with cryptographic verification
- Automatic negotiation and optimization

8.5 Cross-Chain & Interoperability Services

AVM supports bridges between intent graphs and other blockchains or AI services. An intent to "swap tokens on Polygon and bridge to Arbitrum" could be fulfilled by multiple Quarks coordinated across chains and verified in a single PoM block.

This enables:

- Seamless cross-chain operations from a single intent
- Unified verification across heterogeneous systems
- Automatic optimization of cross-chain transactions

9. Future Work & Research Directions

9.1 Scalability and Throughput Enhancements

While AVM's architecture provides significant scalability advantages, several research directions could further improve throughput:

1. **Advanced Parallelism Techniques:** Further research into dynamic load balancing and adaptive parallelism could optimize Quark execution in heterogeneous computing environments.
2. **Sharding of Intent Workflows:** Investigating the potential for workflow sharding, where independent subgraphs of intent executions (DeepFlows) are processed in parallel across different validator sets.
3. **State Compression:** Developing advanced methods for compressing state deltas (δ) and mutation history to reduce storage requirements without compromising verifiability.

9.2 Enhanced Privacy Models

AVM's use of zero-knowledge proofs provides a foundation for privacy, but additional work is needed:

Algorithm 8: Private Intent Execution

Input: Intent I , user private data P

Output: Public mutation μ_{pub} , private proof π_{priv}

1. function ExecutePrivateIntent(I , P)
2. // Generate encrypted DeepFlow
3. $D_{\text{enc}} \leftarrow \text{EncryptDeepFlow}(\text{DecomposeIntent}(I), P)$
- 4.
5. // Execute privately using MPC or FHE
6. $\text{results}_{\text{enc}} \leftarrow \text{SecureCompute}(D_{\text{enc}})$

```

7.
8.    // Generate public mutation with private data commitments
9.     $\mu_{\text{pub}} \leftarrow \text{CreatePublicMutation}(\text{results\_enc})$ 
10.
11.   // Generate private proof viewable only by authorized parties
12.    $\pi_{\text{priv}} \leftarrow \text{GeneratePrivateProof}(D_{\text{enc}}, \text{results\_enc}, P)$ 
13.
14.   return ( $\mu_{\text{pub}}$ ,  $\pi_{\text{priv}}$ )
15. end function

```

Research areas include:

1. **Homomorphic Encryption Integration:** Exploring fully homomorphic encryption (FHE) for private Quark execution
2. **Multi-Party Computation Protocols:** Developing secure multi-party computation (SMPC) techniques for collaborative intent execution
3. **Private Information Retrieval:** Implementing efficient PIR schemes for privacy-preserving data access within DeepFlows

9.3 Quantum-Resistant Advancements

While AVM already incorporates quantum-resistant primitives, ongoing research will focus on:

1. **Lattice-Based Proof Systems:** Developing more efficient lattice-based zero-knowledge proof systems that maintain quantum resistance
2. **Hybrid Cryptographic Approaches:** Combining different post-quantum techniques to optimize for both security and performance
3. **Quantum Random Oracle Models:** Formal security proofs in the quantum random oracle model (QROM)

9.4 Formal Verification and Protocol Analysis

To ensure correctness of the AVM system:

1. **Formal Verification of Core Protocols:** Developing machine-checkable proofs of the PoM consensus protocol and intent execution semantics
2. **Automated Security Analysis:** Creating tools for automated analysis of DeepFlows and mutation conflicts
3. **Compositional Verification:** Developing techniques to verify properties of composed Quarks and mutations

10. Conclusion

The Aintent Virtual Machine represents a fundamental shift in decentralized computation, moving from function-centric execution to intent-driven workflows with cryptographic guarantees. By introducing Deep Intents, DeepFlows, Quarks, and Mutations as primitive constructs, AVM provides a robust foundation for autonomous systems, decentralized applications, and AI-driven workflows.

The integration of zero-knowledge proofs with the novel Proof of Mutations consensus mechanism creates a unique combination of verifiability, scalability, and security—including resistance to quantum attacks through carefully selected cryptographic primitives.

As research continues in areas like advanced parallelism, privacy-preserving computation, and formal verification, AVM has the potential to unlock new paradigms of human-machine interaction and trusted autonomous systems.

References

1. Goldwasser, S., Micali, S., & Rackoff, C. (1989). The knowledge complexity of interactive proof systems. *SIAM Journal on Computing*, 18(1), 186-208.
2. Ben-Sasson, E., Bentov, I., Horesh, Y., & Riabzev, M. (2018). Scalable, transparent, and post-quantum secure computational integrity. *IACR Cryptol. ePrint Arch.*, 2018, 46.
3. Bunz, B., Bootle, J., Boneh, D., Poelstra, A., Wuille, P., & Maxwell, G. (2018). Bulletproofs: Short proofs for confidential transactions and more. In *2018 IEEE Symposium on Security and Privacy (SP)* (pp. 315-334). IEEE.
4. Bernstein, D. J., & Lange, T. (2017). Post-quantum cryptography. *Nature*, 549(7671), 188-194.
5. Bernstein, D. J., Dobraunig, C., Eichlseder, M., Fluhrer, S., Gazdag, S. L., Hülsing, A., ... & Schwabe, P. (2019). SPHINCS+: Submission to the NIST post-quantum project.
6. Ducas, L., Kiltz, E., Lepoint, T., Lyubashevsky, V., Schwabe, P., Seiler, G., & Stehlé, D. (2018). CRYSTALS-Dilithium: A lattice-based digital signature scheme. *IACR Transactions on Cryptographic Hardware and Embedded Systems*, 2018(1), 238-268.
7. Bos, J. W., Ducas, L., Kiltz, E., Lepoint, T., Lyubashevsky, V., Schanck, J. M., ... & Stehlé, D. (2018). CRYSTALS-Kyber: a CCA-secure module-lattice-based KEM. In *2018 IEEE European Symposium on Security and Privacy (EuroS&P)* (pp. 353-367). IEEE.
8. Chiesa, A., Hu, Y., Maller, M., Mishra, P., Vesely, N., & Ward, N. (2020). Marlin: Preprocessing zkSNARKs with universal and updatable SRS. In *Annual International Conference on the Theory and Applications of Cryptographic Techniques* (pp. 738-768). Springer.
9. Buterin, V. (2017). On sharding blockchains. *Ethereum Wiki*.

10. Lamport, L., Shostak, R., & Pease, M. (1982). The Byzantine generals problem. ACM Transactions on Programming Languages and Systems (TOPLAS), 4(3), 382-401.