

# Intention Is All You Need

**Abhimanyu Kashyap**  
aby@helloaxes.com

## **Abstract**

This paper introduces Deep Intent Architecture, a novel framework for decomposing, orchestrating, and executing complex user intents through verifiable workflows. Building on principles of intent engineering, we propose a multi-layered system that translates natural language requests into formalized intent structures, generates optimized execution pathways (DeepFlows), and executes them with cryptographic verification. Our architecture employs a specialized Deep Intent Markup Language (DIML) for structured representation of complex workflows and a novel consensus mechanism called Proof of Mutations for ensuring state synchronization across nodes. Experimental results demonstrate significant improvements in intent fulfillment accuracy, computational efficiency, and verifiability compared to traditional approaches. The proposed architecture provides a foundation for the next generation of agentic systems capable of handling increasingly complex user requests while maintaining transparency, efficiency, and verifiability.

## **1 Introduction**

As artificial intelligence systems advance, there is growing demand for agents capable of understanding and executing complex user intentions with minimal friction. While recent developments in large language models (LLMs) have significantly improved natural language understanding, there remains a substantial gap between understanding intentions and reliably executing them through verifiable pathways. This challenge is particularly evident in systems that must interact with multiple services, allocate computational resources dynamically, and provide users with transparent verification of task completion.

Traditional approaches to intent fulfillment often rely on predefined templates or simplistic decomposition that fails to capture the nuanced relationships between subtasks. Furthermore, the execution of these intents frequently lacks verifiability, making it difficult for users to trust the outcomes or inspect the execution process.

In this paper, we introduce Deep Intent Architecture, a comprehensive framework for intent-driven workflow execution inspired by the transformer architecture's attention mechanisms [1]. Just as transformers revolutionized natural language processing by focusing on attention between input elements, our system focuses on the relationships between intent components to generate optimized execution pathways.

Our contributions include:

- A multi-layered intent decomposition system that transforms natural language requests into formalized intent structures.
- Deep Intent Markup Language (DIML), a specialized language for representing complex workflows.
- DeepFlows, a mechanism for generating and optimizing execution pathways for intent fulfillment.
- Proof of Mutations, a novel consensus algorithm for ensuring state synchronization across nodes.
- A comprehensive architecture for intent execution with cryptographic verification at each step.

## **2 Related Work**

### **2.1 Intent-Based Computing**

Intent-based computing has evolved from simple rule-based systems to more sophisticated approaches using machine learning. Early work by Cohen et al. [2] explored assistant systems that could understand basic user intentions, while more recent work has leveraged LLMs for intent classification [3].

### **2.2 Workflow Orchestration**

Traditional workflow orchestration systems like Apache Airflow [4] and Luigi [5] provide frameworks for defining computational graphs but lack the flexibility to handle dynamic intent decomposition. More recent systems like Temporal [6] improve fault tolerance but still require explicit workflow definition.

### **2.3 Blockchain-Based Verification**

Verifiable computation on blockchain networks has seen significant development, from Ethereum's smart contracts [7] to specialized zero-knowledge proof systems [8]. However, these approaches have not been extensively applied to intent-based workflow execution.

### **2.4 Transformer Architectures**

The transformer architecture introduced in "Attention is All You Need" [1] revolutionized natural language processing by effectively capturing relationships between input elements through attention mechanisms. Our work applies similar principles to intent decomposition and execution.

## **3 System Architecture Overview**

The Deep Intent Architecture consists of the following interconnected layers:

1. Natural Language Capture: Tokenization, context integration
2. Semantic Translation Layer: LLM-based structuring
3. Intent Decomposition Engine: DAG generation
4. DeepFlow Generator: Strategy enumeration

5. Execution Engine: Solver & Quark allocation
6. Agent Virtual Machine (AVM): Blockchain-based execution
7. Proof of Mutations (PoM): zk-SNARK-driven consensus

### 3.1 Natural Language Capture

The entry point to our system is the Natural Language Capture component, which processes user requests expressed in natural language. Unlike traditional intent recognition systems that map utterances to predefined intents, our approach preserves the richness of the original request for deeper processing.

This component performs:

- Initial tokenization and parsing
- Context preservation from previous interactions
- Preliminary intent markers identification
- Parameter extraction for subsequent processing

### 3.2 Semantic Translation Layer

The Semantic Translation Layer transforms the captured natural language into a structured intent representation using advanced language models. This layer performs:

- Deep semantic analysis of user requests
- Extraction of the main goal and subgoals
- Identification of constraints and dependencies
- Parameter formalization for execution planning

The output of this layer is a structured intent representation that serves as input for the Intent Structure Formalization process. We employ a fine-tuned transformer trained on labeled intent graphs to ensure accurate semantic mapping and contextual preservation.

The translation function maps user input to structured intent:

$$\Phi_{LLM} \theta(x) = LLM\theta(x) \rightarrow I$$

Where  $I = (G, C, P, D)$ , representing Goal, Constraints, Parameters, and Dependencies.

#### 3.2.1 Intent Decomposition Algorithm

Algorithm 1 outlines the intent decomposition process:

```
function decompose_intent(raw_intent, context) {
    // Step 1: Parse and extract primary components
    parsed_intent = initial_parse(raw_intent, context)

    // Step 2: Extract the main objective of the intent
    main_goal =
    parsed_intent.match(/\b(goal|objective|task|aim|purpose|intent):?\s*([^\s;,\]+)\/
    i)[2]
```

```

// Step 3: Extract global constraints
constraints =
parsed_intent.scan(/\b(constraint|limitation|restriction|requirement):?\s*([^.
;,+]+)/i).map(match => match[2])

// Step 4: Initialize subgoals collection
subgoals = []

// Step 5: Process remaining content to identify subgoals
remaining_content = parsed_intent

// Step 6: Iteratively extract subgoals
while
(remaining_content.match(/\b(subtask|step|action|operation|phase):?\s*([^.
;,+]+)/i)) {
  // Extract next subgoal
  subgoal_match =
remaining_content.match(/\b(subtask|step|action|operation|phase):?\s*([^.
;,+]+)/i)
  next_subgoal = subgoal_match[2]

  // Find subgoal-specific constraints
  subgoal_constraints = remaining_content.scan(new
RegExp(`\b(for|with|using|${next_subgoal})\s+requires)\s+([^.
;,+])`,
'i')).map(match => match[2])

  // Identify dependencies between subgoals
  dependency_matches = remaining_content.scan(new
RegExp(`\b(after|following|dependent
on|requires)\s+([^.
;,+])\s+${next_subgoal}`, 'i'))
  subgoal_dependencies = dependency_matches.map(match =>
    subgoals.findIndex(sg => sg.objective.includes(match[2]))
  ).filter(idx => idx >= 0).map(idx => subgoals[idx].id)

  // Add structured subgoal to collection
  subgoals.push({
    "id": `SG-${subgoals.length + 1}`,
    "objective": next_subgoal,
    "constraints": subgoal_constraints,
    "dependencies": subgoal_dependencies
  })

  // Update remaining content to remove processed subgoal

```

```

    remaining_content = remaining_content.replace(subgoal_match[0], '')
}

// Step 7: Assemble full intent structure
intent_structure = {
  "intentId": generateHash(main_goal + Date.now()),
  "rawIntent": raw_intent,
  "mainGoal": {
    "objective": main_goal,
    "constraints": constraints,
    "priority": determine_priority(main_goal, context)
  },
  "subGoals": subgoals
}

return intent_structure
}

```

This generates DAG-based task graphs for flexible workflow planning.

### 3.3 Intent Structure Formalization

Once decomposed, the intent must be formalized into a structure that can guide execution planning. Our approach uses a specialized representation that captures the hierarchical nature of intents and their constraints.

The formal intent structure includes:

- A unique identifier for the intent
- The original raw intent text
- A main goal with associated constraints and priority
- A set of subgoals, each with its own constraints and dependencies

Example of a formalized intent structure:

```

{
  "intentId": "travel-planning-nyc-apr-2025",
  "rawIntent": "I'd like to travel from London to New York in April for 10
days with a",
  "mainGoal": {
    "objective": "Plan NYC trip",
    "constraints": [
      "April 2025",
      "10 days",

```

```
    "budget $5,000"
  ],
  "priority": "high"
},
"subGoals": [
  {
    "id": "SG-1",
    "objective": "Book flights",
    "dependencies": [],
    "constraints": [
      "from London",
      "to New York",
      "April 2025",
      "budget allocation: $"
    ]
  },
  {
    "id": "SG-2",
    "objective": "Reserve accommodation",
    "dependencies": [],
    "constraints": [
      "New York",
      "10 nights",
      "budget allocation: $2,300"
    ]
  },
  {
    "id": "SG-3",
    "objective": "Plan activities",
    "dependencies": [
      "SG-1",
      "SG-2"
    ],
    "constraints": [
      "budget allocation: $700"
    ]
  }
]
}
```

## 4 DeepFlow Generation and Optimization

With the formalized intent structure, the system generates multiple potential execution pathways called DeepFlows. Each DeepFlow represents a different approach to fulfilling the user's intent.

We define candidate DeepFlows as:

$$D_i = (S_i, R_i, T_i)$$

Where:

$S_i$ : Sequence of Steps (DAG)

$R_i$ : Resource cost

$T_i$ : Estimated time

The selected path satisfies:  $D^* = \arg \min D_i [\alpha T_i + \beta R_i - \gamma V_i]$  with verifiability score  $V_i$  derived from zk-proof metrics.

Where:

$T_i$ : Execution time

$R_i$ : Resource cost

$V_i$ : zk-verifiability score

DeepFlow Scoring Algorithm is as follows :  $D^* = \arg \min D_i [\alpha T_i + \beta R_i - \gamma V_i]$

```
function score_deepflow(flow, parameters) {
  # Extract scoring parameters
  const  $\alpha$  = parameters.time_weight || 1.0
  const  $\beta$  = parameters.resource_weight || 1.0
  const  $\gamma$  = parameters.verifiability_weight || 1.0

  # Calculate execution time score (normalized)
  const T = flow.metadata.match(/EstimatedDuration:(\d+(\.\d+)?) [sm]/) [1] /
parameters.max_acceptable_time

  # Calculate resource usage score (normalized)
  const resource_match = flow.metadata.match(/EstimatedCost:(\d+(\.\d+)?) /i)
  const R = resource_match ? parseFloat(resource_match[1]) /
parameters.max_acceptable_cost : 1.0

  # Calculate verifiability score based on proof coverage
  const steps = flow.match(/<Step[^>]*>/g) || []
  const verified_steps = flow.match(/<Verification[^>]*>/g) || []
}
```

```

const V = verified_steps.length / Math.max(steps.length, 1)

# Apply scoring formula with weights
return ( $\alpha$  * T) + ( $\beta$  * R) - ( $\gamma$  * V)
}

function select_optimal_deepflow(candidate_flows, parameters) {
  let best_score = Infinity
  let optimal_flow = null

  for (const flow of candidate_flows) {
    const score = score_deepflow(flow, parameters)
    if (score < best_score) {
      best_score = score
      optimal_flow = flow
    }
  }

  return optimal_flow
}

```

#### 4.1 LLM-Driven Intent Orchestrator

The Intent Orchestrator uses the decomposed intent to generate multiple solution pathways. These pathways are structured using the Deep Intent Markup Language (DIML), which creates a directed acyclic graph (DAG) of execution steps.

#### 4.2 Deep Intent Markup Language (DIML)

DIML is a specialized markup language that provides a standardized format for representing complex workflows. It supports both human readability and machine execution, enabling visualization of intent execution pathways.

DIML enables standardization across workflow steps. A DeepFlow consists of:

- Metadata (cost, duration, authorship)
- Step-wise execution DAGs
- Bound variables and resource references

DIML provides both machine readability and developer accessibility for debugging and reuse.

Example of a DeepFlow expressed in DIML:

```

<DeepFlow id="DF-Travel-NYC-001" intent="travel-planning-nyc-apr2025">

  <Metadata>

    <Author>Intent Orchestrator</Author>

```



```
<Created>2025-04-03T10:15:32Z</Created>

<EstimatedCost>0.15</EstimatedCost>

<EstimatedDuration>45s</EstimatedDuration>

</Metadata>

<Resources>

  <Resource id="R1" type="FlightAPI" provider="GlobalFlights" />

  <Resource id="R2" type="AccommodationAPI" provider="StayFinder" />

  <Resource id="R3" type="ActivityPlanner" provider="TripMaster" />

</Resources>

<Steps>

  <Step id="S1" depends="">

    <Action resource="R1">

      <Query>

        findFlights(origin="London", destination="New York",

          departDate="2025-04-22", returnDate="2025-05-02",

          maxBudget=2000)

      </Query>

      <Output bind="flightOptions" />

    </Action>

  </Step>

  <Step id="S2" depends="">

    <Action resource="R2">

      <Query>

        findAccommodation(location="New York", neighborhood="Midtown",

          checkIn="2025-04-22", checkOut="2025-05-02",

          maxBudget=2300)

      </Query>
```

```

        <Output bind="accommodationOptions" />

    </Action>

</Step>

<Step id="S3" depends="S1,S2">

    <Action resource="R3">

        <Query>

            planActivities(location="New York",

                arrival=${flightOptions.arrivalTime},

                departure=${flightOptions.departureTime},

                hotel=${accommodationOptions.location},

                interests=["museums", "Broadway", "food"],

                budget=700)

        </Query>

        <Output bind="activities" />

    </Action>

</Step>

</Steps>

<Output>

    <Combine>

        <Item source="flightOptions" />

        <Item source="accommodationOptions" />

        <Item source="activities" />

    </Combine>

</Output>

</DeepFlow>

```

### 4.3 User Decision Point

After DeepFlows are generated, users can review multiple workflow options and select or modify the approach that best aligns with their preferences.

## 5 Execution Engine and Verifiability

Once a DeepFlow is selected, the Execution Engine processes it through specialized solvers and resource allocation.

A selected DeepFlow is executed by:

- Assigning solvers
- Allocating quarks
- Generating step-by-step zero-knowledge proofs

Every state mutation must pass:

- $\mu_i: S \rightarrow S'$
- $\text{Verify}(\pi_i, \mu_i, S) = \text{True}$

Execution runtime includes:

- Solvers: Task-type specific agents
- Quarks: Atomic compute units
- Trace Logs & Proofs: zk-SNARKs generated per mutation

### 5.1 Solver Allocation

The system assigns appropriate solvers for each step in the workflow. Solvers are matched to the specific requirements of each task, and resource arrays are allocated for computation.

### 5.2 Quark Utilization

Computational resources (Quarks) are assigned to specific tasks based on their requirements. Resource usage is monitored and optimized in real-time, with Quark allocation adjusting dynamically based on execution needs.

```
function allocate_resources(step, resource_pool) {  
  
  // Step 1: Analyze step complexity  
  
  const complexity_pattern = /<Complexity[^>]*value="([^\"]+)"[^>]*\>/i  
  
  const complexity_match = step.match(complexity_pattern)  
  
  const complexity = complexity_match ? parseFloat(complexity_match[1]) : 1.0  
  
  
  // Step 2: Determine resource requirements based on complexity and step  
  type  
  
  const required_resources = {}  
}
```

```

// Extract operation type from step

const operation_pattern = /<Operation[^>]*type="([^\"]+)"[^>]*\>/i

const operation_match = step.match(operation_pattern)

const operation_type = operation_match ? operation_match[1] : "default"

// Calculate resources by type

if (operation_type.match(/computation|processing|analysis/i)) {

    required_resources["compute"] = Math.ceil(complexity * 2.5)

    required_resources["memory"] = Math.ceil(complexity * 1.8)

} else if (operation_type.match(/storage|database|retrieval/i)) {

    required_resources["storage"] = Math.ceil(complexity * 3.0)

    required_resources["io"] = Math.ceil(complexity * 1.5)

} else if (operation_type.match(/api|external|service/i)) {

    required_resources["network"] = Math.ceil(complexity * 2.0)

    required_resources["io"] = Math.ceil(complexity * 1.0)

} else {

    // Default allocation for unknown operations

    required_resources["compute"] = Math.ceil(complexity * 1.0)

    required_resources["memory"] = Math.ceil(complexity * 1.0)

    required_resources["network"] = Math.ceil(complexity * 0.5)

}

// Step 3: Select optimal solver based on operation type and requirements

const solver_id = find_best_solver(step.resource_type, required_resources)

// Step 4: Allocate quarks from resource pool

const allocated_quarks = {}

for (const [resource_type, amount] of Object.entries(required_resources)) {

    // Get available quarks of this type

```

```

const available_quarks = resource_pool.filter(quark =>
    quark.type === resource_type && quark.status === "available"
)

// Allocate up to the required amount

const allocation_count = Math.min(amount, available_quarks.length)

    allocated_quarks[resource_type] = available_quarks.slice(0,
allocation_count).map(q => q.id)

    // Mark allocated quarks as busy

    for (const quark_id of allocated_quarks[resource_type]) {

        resource_pool.find(q => q.id === quark_id).status = "allocated"

    }

}

return {

    solver: solver_id,

    quarks: allocated_quarks

}

}

```

### 5.3 Proof Generation

As tasks are completed, cryptographic proofs are generated to verify the correct execution of each workflow step. The complete chain of proofs ensures the integrity of the entire process.

## 6. Agent Virtual Machine and Consensus

To enable decentralized execution and verification of intents, we implement an Agent Virtual Machine (AVM) as a Layer 1 blockchain. This blockchain maintains the state of accounts, agentic activity, and intents as transactions.

### 6.1 Blockchain Architecture

The AVM blockchain architecture consists of:

- Bootnode: Serves as a batcher, sequencer, and proposer
- Light clients: Access the network with minimal resources

- Peer nodes: Facilitate network communication
- Block structure: Contains sequenced intent transactions

Each block contains:

$$B = \{\mu_1, \mu_2, \dots, \mu_n\}$$

And each mutation is cryptographically proven:

$$\forall \mu_i, \exists \pi_i: \text{ZK-Verify}(\mu_i, \pi_i, S_{i-1}) = \text{True}$$

## 6.2 Proof of Mutations Consensus Algorithm

Our novel consensus mechanism, Proof of Mutations (PoM), ensures that all state changes (mutations) are synchronized across nodes. Unlike Proof of Work or Proof of Stake, PoM focuses on validating the correctness of state transitions rather than resource expenditure or stake.

PoM ensures near-instant finality with block times of 1-2 seconds, enabling responsive intent execution while maintaining consensus across the network.

Consensus is reached when:

- All mutations are valid
- The final state hash is agreed upon
- The proof chain is continuous

PoM avoids the energy cost of PoW and the centralization of PoS, offering verifiability at each computational step.

```
function validate_mutation(current_state, mutation, proofs) {
  // Step 1: Verify the mutation signature

  const signature_pattern = /<Signature>([A-Za-f0-9]+)<\|Signature>/

  const signature_match = mutation.match(signature_pattern)

  if (!signature_match || !verify_cryptographic_signature(
    signature_match[1],
    mutation.replace(signature_pattern, ''),
    mutation.match(/<Author>([^\|]+)<\|Author>/) [1]
  )) {
```

```

    return { valid: false, reason: "Invalid signature" }

}

// Step 2: Verify mutation applicability to current state

const state_version =
current_state.match(/<StateVersion>(\d+)<\/StateVersion>/)[1]

const mutation_base_version =
mutation.match(/<BasedOnVersion>(\d+)<\/BasedOnVersion>/)[1]

if (parseInt(state_version) !== parseInt(mutation_base_version)) {

    return { valid: false, reason: "Mutation cannot be applied to current
state version" }

}

// Step 3: Verify each execution step has valid proof

const steps = mutation.match(/<ExecutionStep id="([^\"]+)"[^\>]*>/g) || []

for (const step_tag of steps) {

    const step_id = step_tag.match(/id="([^\"]+)"/)[1]

    const proof_id = step_tag.match(/proof="([^\"]+)"/)[1]

    if (!proofs[proof_id] || !verify_step_proof(

        extract_step_details(mutation, step_id),

        proofs[proof_id]

    )) {

        return { valid: false, reason: `Invalid proof for step ${step_id}` }

    }

}

}

// Step 4: Apply mutation to generate new state and verify hash

const new_state = apply_mutations_to_state(current_state, mutation)

```

```

const expected_hash =
mutation.match(/<ExpectedStateHash>([A-Za-f0-9]+)<\/ExpectedStateHash>/)[1]

const actual_hash = calculate_hash(new_state)

if (expected_hash !== actual_hash) {

  return {

    valid: false,

    reason: `State hash mismatch. Expected: ${expected_hash}, Got:
${actual_hash}`

  }

}

return { valid: true, new_state }

}

function verify_step_proof(step, proof) {

  // Implementation of zero-knowledge proof verification

  const input_hash = calculate_hash(step.inputs)

  const output_hash = calculate_hash(step.outputs)

  const operation_hash = calculate_hash(step.operation)

  // Verify the proof validates this operation produced these outputs from
these inputs

  return zk_verify(

    proof,

    { input_hash, output_hash, operation_hash }

  )

}

```



### 6.3 Verifiable Services

Intent solvers claiming to process DeepFlows must submit micro-transactions within the main intent transaction to external provers. This creates a robust proof chain of the activities conducted by the execution engine.

The verification process includes:

1. Pre-execution attestation of solver capabilities
2. Step-by-step proof generation during execution
3. Post-execution verification of the complete proof chain
4. State mutation validation across the network

We evaluated our Deep Intent Architecture against traditional intent processing systems using several metrics:

Metric	Definition
Fulfillment Accuracy	Total Intents / Valid Intents Executed
Resource Efficiency	Optimal Quark Budget / Actual Quark Usage
Verifiability Score	Total Execution Steps / Verified Steps $\in [0, 1]$
Execution Time	Time from intent expression to fulfillment

Table 1 presents the results across three intent complexity levels:

Metric	Simple Intents		Complex Intents		Multi-Stage Intents	
	Traditional	Deep Intent	Traditional	Deep Intent	Traditional	Deep Intent
Fulfillment Accuracy (%)	92.3	94.5	78.6	91.2	63.4	87.9
Resource Usage (norm.)	1.0	0.87	1.0	0.74	1.0	0.68
Verification Score (0-1)	0.65	0.98	0.58	0.97	0.42	0.96
Execution Time (s)	12.3	9.8	34.7	18.5	86.5	31.2

Our architecture shows significant improvements across all metrics, with the most pronounced benefits for complex and multi-stage intents.

## **8. Case Study: Complex Travel Planning**

To demonstrate the capabilities of our architecture, we present a case study of complex travel planning intent execution. The user intent was:

"I'd like to travel from London to New York in April for 10 days. I want a hotel near the subway for easy transportation. My budget is \$5,000 for flights, accommodation, and food. Could you help me plan everything and make the bookings?"

The system processed this intent through the following stages:

1. Natural Language Capture: Extracted initial parameters (origin: London, destination: New York, timeframe: April, duration: 10 days, budget: \$5,000).
2. Additional information gathering: System requested and received details about date preferences, departure airport, accommodation preferences, and activity interests
3. Semantic Translation Layer: Converted detailed parameters into a formal intent structure with constraints and feasibility analysis.
4. DeepFlow Generation: Created multiple strategic approaches (luxury-focused, balanced, experience-maximizing).
5. User Selection: User selected the balanced approach with Virgin Atlantic flights.
6. Execution Engine: Processed bookings for flights and accommodation while maintaining budget constraints.
7. Activity Planning: Generated an itinerary with museums, Broadway shows, and authentic dining experiences.
8. Outcome Delivery: Provided comprehensive travel package with all confirmations and details.

The entire process was executed with cryptographic verification at each step, ensuring the user could trust the bookings and arrangements made on their behalf.

## **9. Discussion and Future Work**

The Deep Intent Architecture represents a significant advancement in intent-based systems, providing a framework for decomposing, planning, and executing complex user requests with verifiable outcomes. Our results demonstrate particular advantages in handling multi-stage intents that require coordination across multiple services and time periods.

Several directions for future work emerge from our research:

1. Improving the intent decomposition system to handle more ambiguous requests.
2. Enhancing the DeepFlow optimization process to better balance efficiency and user preferences.
3. Extending the verification system to handle privacy-preserving proofs.
4. Developing specialized solvers for domain-specific intent execution.
5. Integrating the system with existing web services and APIs.

## 10. Conclusion

In this paper, we introduced Deep Intent Architecture, a comprehensive framework for intent-driven workflow execution. Our approach transforms natural language requests into formalized intent structures, generates optimized execution pathways, and executes them with cryptographic verification. The architecture's blockchain-based Agent Virtual Machine ensures transparent and verifiable intent fulfillment across a decentralized network.

Just as the transformer's attention mechanism revolutionized natural language processing by focusing on relationships between input elements, our Deep Intent Architecture focuses on the relationships between intent components to generate optimized execution pathways. Our experimental results demonstrate significant improvements in intent fulfillment accuracy, computational efficiency, and verifiability compared to traditional approaches.

By providing a robust foundation for intent decomposition, orchestration, and verified execution, our work opens new possibilities for agentic systems that can reliably fulfill increasingly complex user requests while maintaining transparency and trust.

## References

- [1] Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., & Polosukhin, I. (2017). Attention is all you need. In *Advances in neural information processing systems* (pp. 5998- 6008).
- [2] Cohen, P. R., Cheyer, A., Wang, M., & Baeg, S. C. (1994). An open agent architecture. In *Proceedings of the AAAI Spring Symposium on Software Agents* (pp. 1-8).
- [3] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J. D., Dhariwal, P., ... & Amodei, D. (2020). Language models are few-shot learners. *Advances in neural information processing systems*, 33, 1877-1901.
- [4] Apache Software Foundation. (2018). Apache Airflow. <https://airflow.apache.org/>
- [5] Spotify. (2015). Luigi. <https://github.com/spotify/luigi>
- [6] Temporal Technologies, Inc. (2020). Temporal workflow engine. <https://temporal.io/>
- [7] Buterin, V. (2014). Ethereum: A next-generation smart contract and decentralized application platform. White paper.
- [8] Ben-Sasson, E., Chiesa, A., Garman, C., Green, M., Miers, I., Tromer, E., & Virza, M. (2014). Zerocash: Decentralized anonymous payments from bitcoin. In *IEEE Symposium on Security and Privacy* (pp. 459-474).
- [9] Aintent Protocol. (2025). Enhanced Documentation for Deep Intent Architecture.
- [10] Deep Intent System. (2025). Detailed User Interaction Flow for Travel Planning.